Text generation

How do we generate sentences?

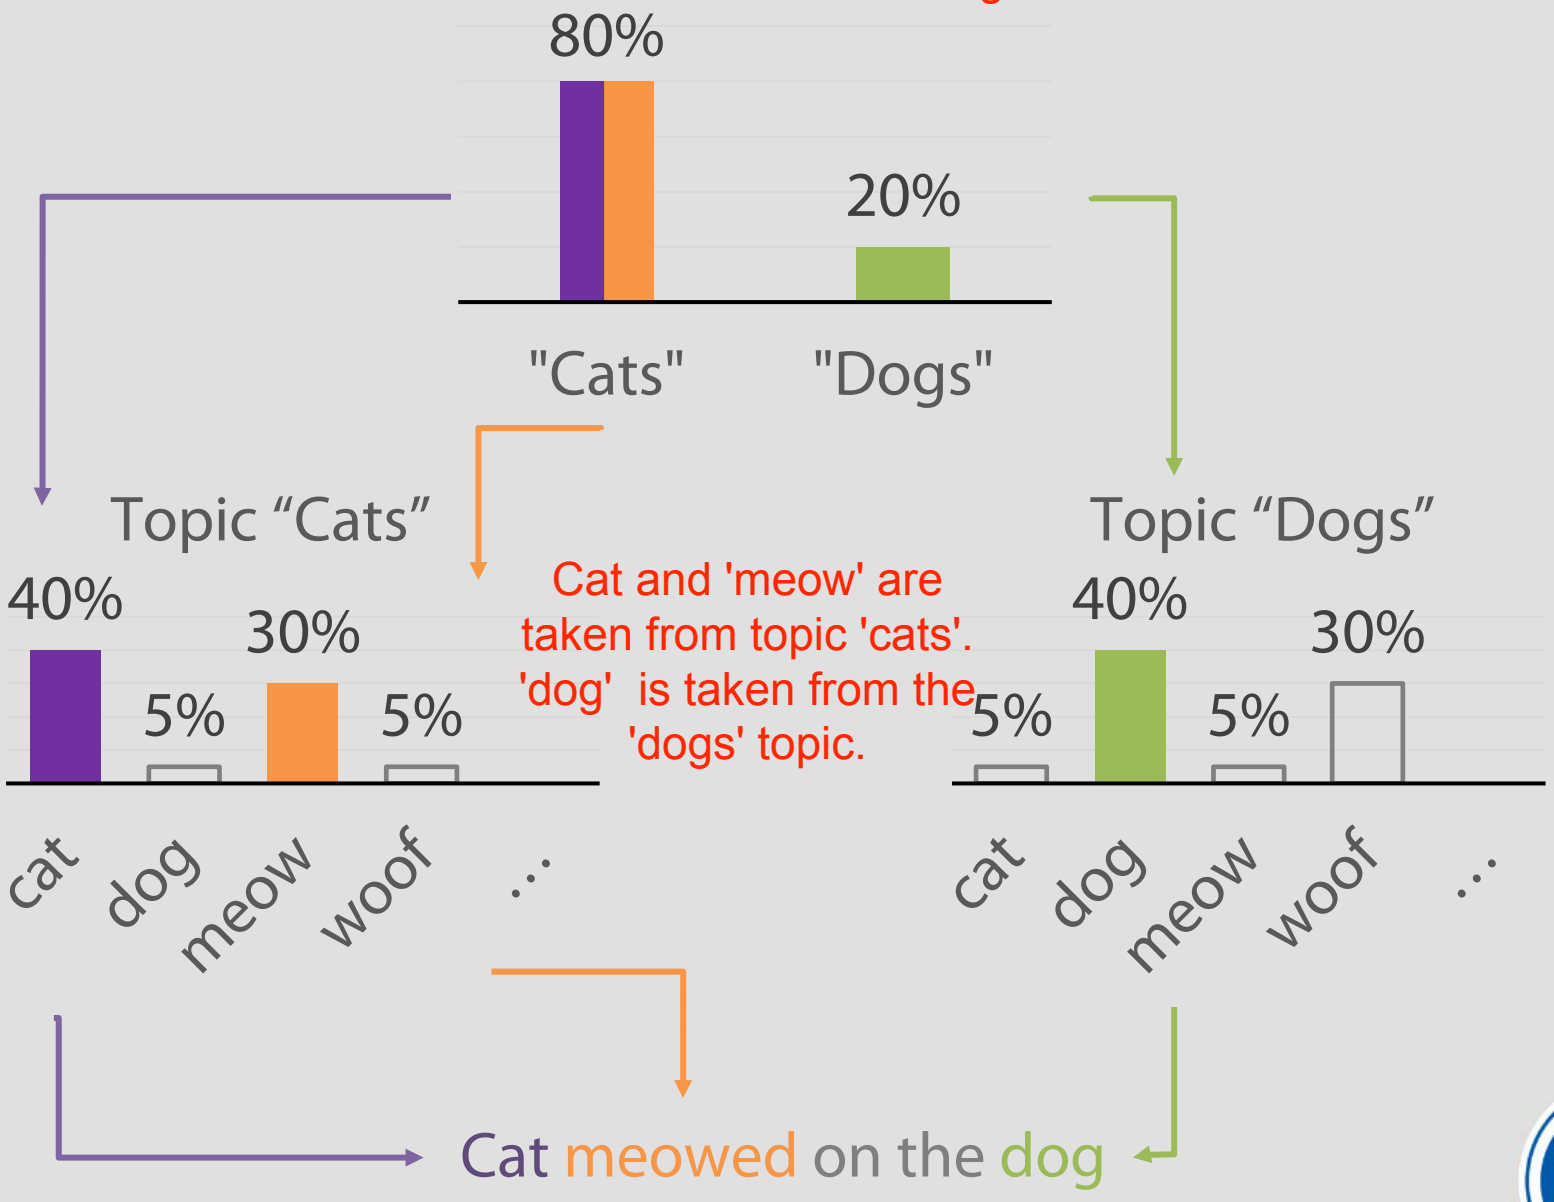

Model

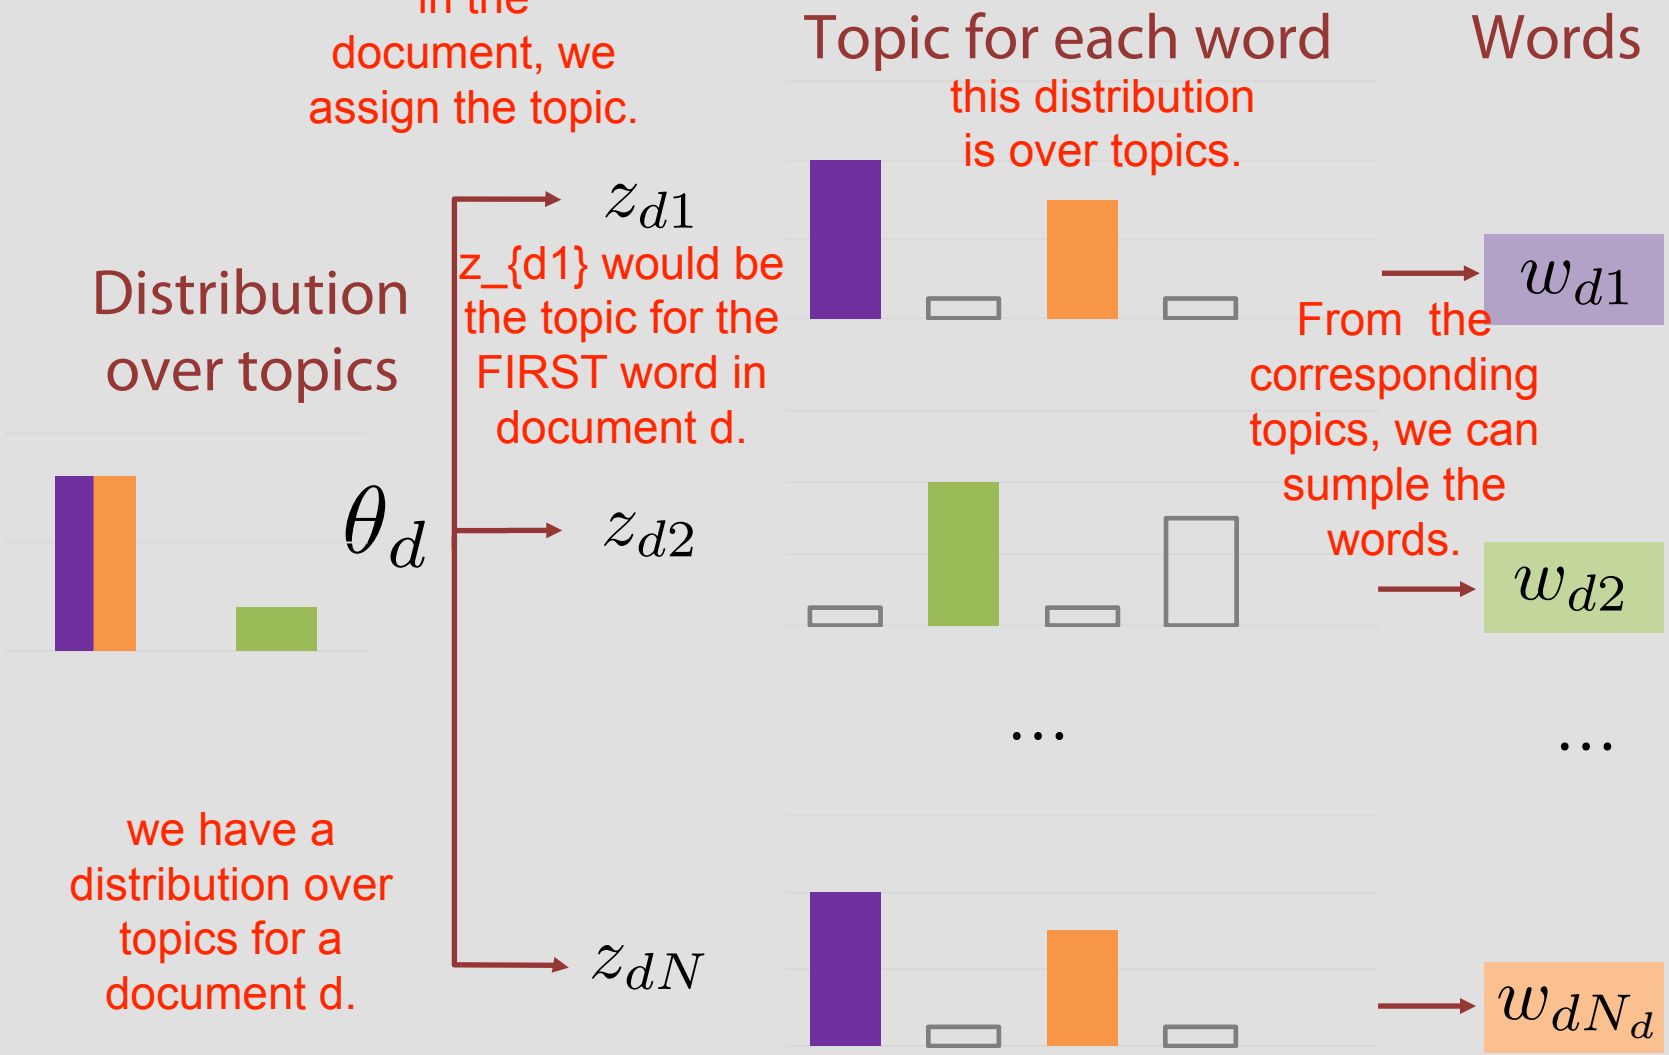

$$z_{dn} \in \{1..T\} \quad w_{dn} \in \{1..V\}$$

There are T topics.

There are V words in the vocabulary.

LDA Model

What we did in the previous slide leads us to this:

A bayesian network in plate notation.

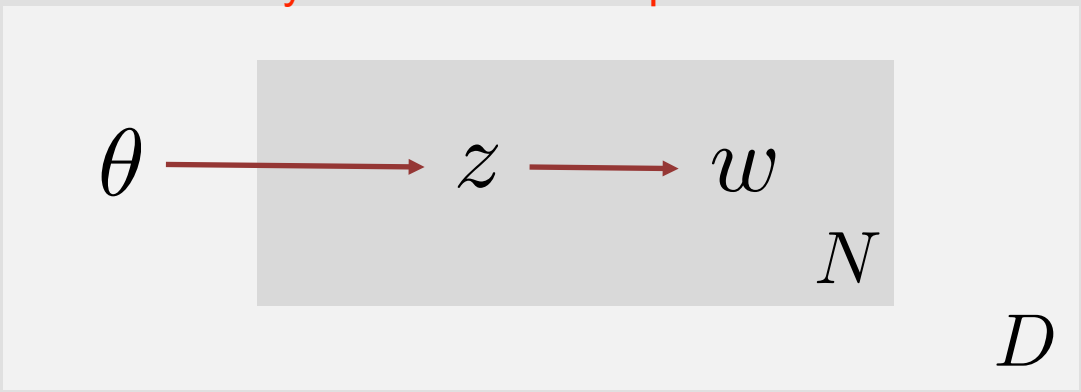

How to interpret plate notation: The theta here is in the D box, which means that we repeat this box for the D documents.
z->w is in the N box. This means that for each document, we have N of these (N words).

in the document,

for each word

select topic

select word from topic

$$p(W, Z, \Theta) = \prod_{d=1}^{D} p(\theta_d) \prod_{n=1}^{N_d} p(z_{dn}|\theta_d) p(w_{dn}|z_{dn})$$

for each document

generate topic probabilities

LDA Model

$$p(W, Z, \Theta) = \prod_{d=1}^{D} p(\theta_d) \prod_{n=1}^{N_d} p(z_{dn}|\theta_d) p(w_{dn}|z_{dn})$$

$$p(\theta_d) \sim \mathrm{Dir}(\alpha)$$

sum of all theta_d sum up to 1. Why?
I suppose the k-simplex model seems pretty 'fitting'. A document can be really heavy on some topic, but may also touch on other topics too.

$$p(z_{dn}|\theta_d) = \theta_{dz_{dn}}$$

$$p(w_{dn}|z_{dn}) = \Phi_{z_{dn} w_{dn}}$$

Row z_dn, column w_dn.

Constraints:

$$\Phi_{tw} \geq 0$$

$$\sum_{w} \Phi_{tw} = 1$$

TODO: WHAT is tw here?

LDA Model

Known: W data

Unknown: Φ parameters, distribution over words for each topic

Unknown: Z latent variables, topic of each word

Unknown: Θ latent variables, distribution over topics for each document